

Logistic Regression

Loading Data and Libraries

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
plt.style.use("ggplot")
%matplotlib inline
```

```
In [2]: from pylab import rcParams
rcParams["figure.figsize"] = 20, 8
```

```
In [3]: df_dmv = pd.read_csv("DMV_Written_Tests.csv")
df_dmv.head()
```

```
Out[3]:
```

	DMV_Test_1	DMV_Test_2	Results
0	34.623660	78.024693	0
1	30.286711	43.894998	0
2	35.847409	72.902198	0
3	60.182599	86.308552	1
4	79.032736	75.344376	1

```
In [4]: df_dmv.shape
```

```
Out[4]: (100, 3)
```

```
In [5]: test_scores = df_dmv[["DMV_Test_1", "DMV_Test_2"]].values
results = df_dmv["Results"].values
```


Data Visualization

```

In [6]: passed = (results == 1).reshape(100, 1)
failed = (results == 0).reshape(100, 1)

ax = sns.scatterplot(x = test_scores[passed[:, 0], 0],
                    y = test_scores[passed[:, 0], 1],
                    marker = "^",
                    color = "blue", s = 120)
sns.scatterplot(x = test_scores[failed[:, 0], 0],
                y = test_scores[failed[:, 0], 1],
                marker = "x",
                color = "red", s = 120)
ax.set(xlabel = "DMV Written Test 1 scores", ylabel = "DMV Written Test 1 scores")
ax.legend(["Passed", "Failed"])
plt.show()

```



Logistic Sigmoid Function $\sigma(z)$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

```
In [7]: def sigmoid_function(x):
        return 1 / (1 + np.exp(-x))
```

```
In [8]: sigmoid_function(0)
```

```
Out[8]: 0.5
```

Computing the Cost Function $J(\theta)$ and Gradient

The objective of logistic regression is to minimize the cost function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - (h_{\theta}(x^{(i)})))]$$

where the gradient of the cost function is given by

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

```
In [9]: def cost_function(theta, x, y):
        m = len(y)
        y_pred = sigmoid_function(np.dot(x, theta))
        error = (y * np.log(y_pred) + (1 - y) * np.log(1 - y_pred))
        cost = - 1 / m * sum(error)
        gradient = 1 / m * np.dot(x.transpose(), (y_pred - y))
        return cost[0], gradient
```

Cost and Gradient at Initialization

```
In [10]: mean_scores = np.mean(test_scores, axis = 0)
std_scores = np.std(test_scores, axis = 0)
test_scores = (test_scores - mean_scores) / std_scores

rows = test_scores.shape[0]
cols = test_scores.shape[1]

X = np.append(np.ones((rows, 1)), test_scores, axis = 1)
y = results.reshape(rows, 1)

theta_init = np.zeros((cols + 1, 1))
cost, gradient = cost_function(theta_init, X, y)

print(" Cost at Initialization is", cost)
print(" Gradients at Initialization is", gradient)

Cost at Initialization is 0.69314718056
Gradients at Initialization is [[-0.1      ]
 [-0.28122914]
 [-0.25098615]]
```

Gradient Descent

Minimize the cost function $J(\theta)$ by updating the below equation and repeat until convergence $\theta_j := \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$ (simultaneously update θ_j for all j)

```
In [11]: def gradient_descent(X, y, theta, alpha, iterations):
costs = []
for i in range(iterations):
    cost, gradient = cost_function(theta, X, y)
    theta -= (alpha * gradient)
    costs.append(cost)
return theta, costs
```

```
In [12]: theta, costs = gradient_descent(X, y, theta_init, 1, 200)
```

```
In [13]: print("Theta after gradient descent: ", theta)
print("Resulting cost: ", costs[-1])
```

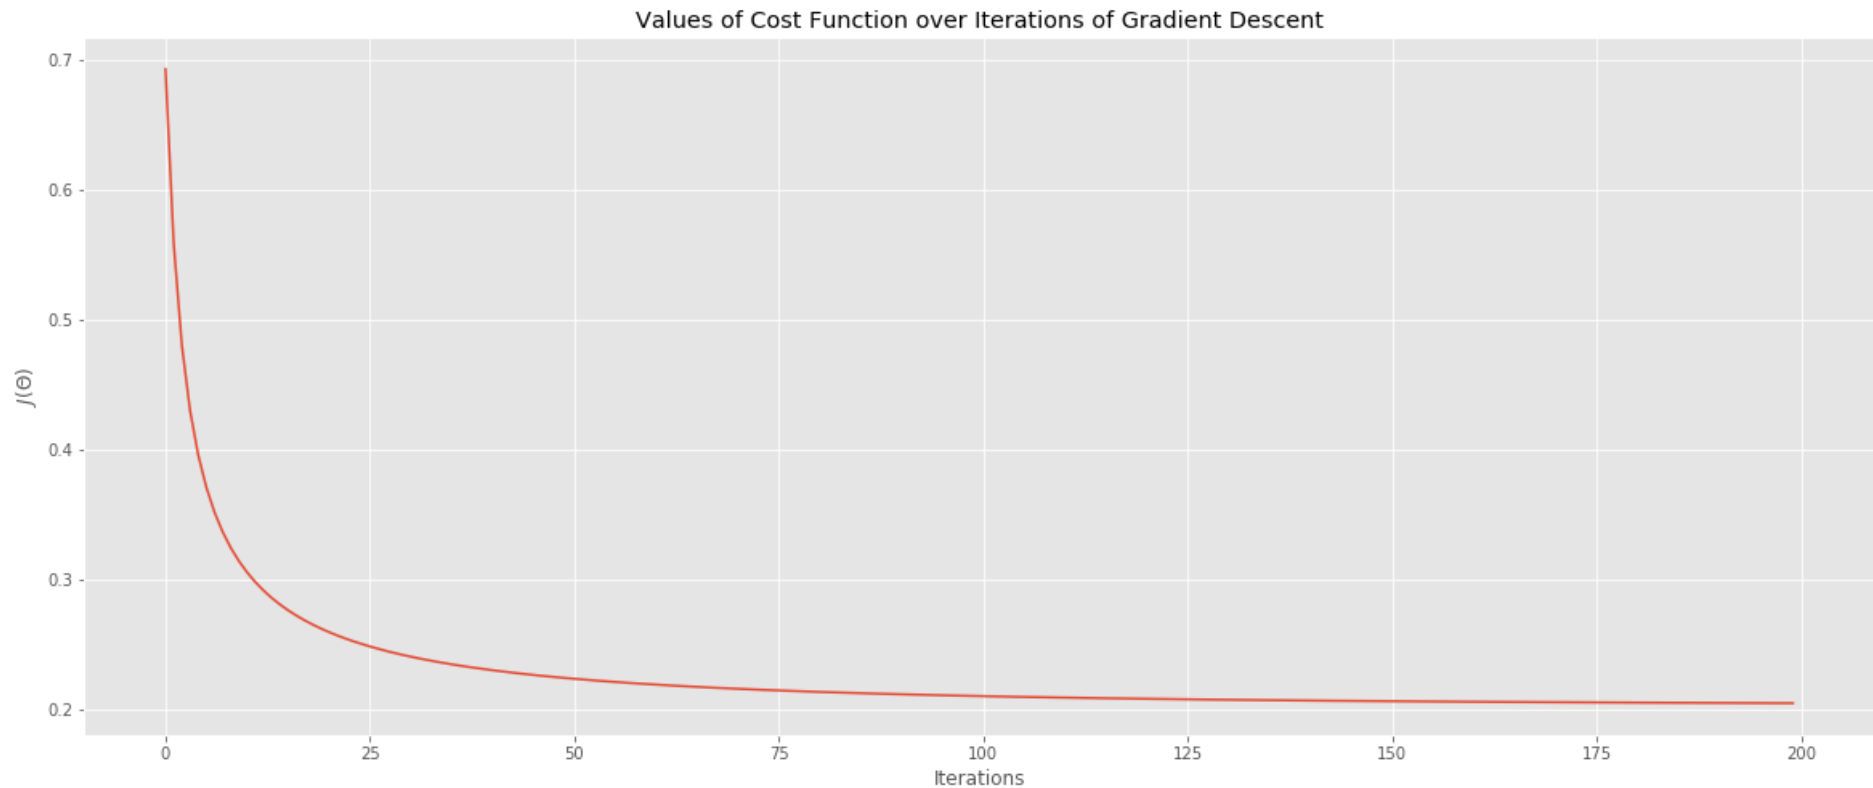
```
Theta after gradient descent: [[ 1.50850586]
 [ 3.5468762 ]
 [ 3.29383709]]
Resulting cost: 0.204893820351
```

Plotting the Convergence of $J(\theta)$

Plot $J(\theta)$ against the number of iterations of gradient descent:

```
In [14]: plt.plot(costs)
plt.xlabel("Iterations")
plt.ylabel("$J(\theta)$")
plt.title("Values of Cost Function over Iterations of Gradient Descent")
```

Out[14]: <matplotlib.text.Text at 0x8cdc110>



Plotting the decision boundary

$h_{\theta}(x) = \sigma(z)$, where σ is the logistic sigmoid function and $z = \theta^T x$

When $h_{\theta}(x) \geq 0.5$ the model predicts class "1":

$$\Rightarrow \sigma(\theta^T x) \geq 0.5$$

$$\Rightarrow \theta^T x \geq 0 \text{ predict class "1"}$$

Hence, $\theta_1 + \theta_2 x_2 + \theta_3 x_3 = 0$ is the equation for the decision boundary, giving us


$$x_3 = \frac{-(\theta_1 + \theta_2 x_2)}{\theta_3}$$

```
In [15]: ax = sns.scatterplot(x = X[passed[:, 0], 1],
                             y = X[passed[:, 0], 2],
                             marker = "^",
                             color = "blue", s = 90)
sns.scatterplot(x = X[failed[:, 0], 1],
                y = X[failed[:, 0], 2],
                marker = "x",
                color = "red", s = 90)

ax.set(xlabel = "DMV Written Test 1 scores", ylabel = "DMV Written Test 1 scores")
ax.legend(["Passed", "Failed"])

x_boundary = np.array([np.min(X[:, 1]), np.max(X[:, 1])])
y_boundary = -(theta[0] + theta[1] * x_boundary) / theta[2]

sns.lineplot(x = x_boundary, y = y_boundary, color = "black")
plt.show()
```



Predictions using the optimized θ values

$$h_{\theta}(x) = x\theta$$

```
In [16]: def predict(theta, x):  
         results = x.dot(theta)  
         return results > 0
```

```
In [17]: p = predict(theta, X)  
print("Training Accuracy: ", sum(p == y)[0], "%")
```

Training Accuracy: 89 %

```
In [18]: test = np.array([60, 79])  
test = (test - mean_scores) / std_scores  
test = np.append(np.ones(1), test)  
probability = sigmoid_function(test.dot(theta))  
  
print("A person who scores 60 and 79 on their DMV written test have a", np.round(probability[0], 2), "probability of passing")
```

A person who scores 60 and 79 on their DMV written test have a 0.94 probability of passing